


EE3731C Pattern Recognition 1

BT Thomas Yeo


ECE, CIRC, Sinapse, Duke-NUS, HMS

Curse of Dimensionality

- Data dimensionality = number of “free” parameters in data / model
 - $1M$ -pixel image $\implies 1M$ dimensions
 - Text document has N words $\implies N$ dimensions
- Dimensionality depends on problem definition
 - Word analysis of text document \implies dimensionality = # words
 - Character analysis of text document \implies dimensionality = # characters
- Curse of dimensionality: pattern recognition harder for big dimensions
 - Joint distribution of binary variables $p(x_1, \dots, x_N)$ has $2^N - 1$ parameters 
 - Need a lot of data to estimate so many parameters

Intrinsic Dimensions


- Given data has D dimensions, hope “intrinsic” data dimensions d less than D
- “Random” images do not look like “real” images \implies image pixels not independent \implies true dimensions less than number of pixels



2-dimensional Points ($D = 2$)
with 1 “intrinsic” dimension ($d = 1$)




“Real” Image




“Random” Image
(Each pixel independent with
uniform $U[0, 1]$ distribution)

Dimensionality Reduction


- Dimensionality reduction: map high-dimensional data into lower dimensions
 - Also known as feature extraction or feature reduction
 - Different assumptions lead to different mapping
- Principal component analysis (PCA) assumes data lies on linear subspace



1-dimensional **linear** subspace
(i.e., can be represented by **straight line**)



2-dimensional **linear** subspace
(i.e., can be represented by **flat plane**)



1-dimensional **nonlinear** space
representable by 1-D **curve**
(PCA cannot handle this)

Linear System Revision

Linear System (or Transformation)

- Suppose system T maps input $x \in \mathbb{R}^n$ to output $y \in \mathbb{R}^m$
 - If system is linear, then $y_1 = T(x_1)$ and $y_2 = T(x_2)$ implies $T(ax_1 + bx_2) = aT(x_1) + bT(x_2) = ay_1 + by_2$
 - All linear systems can be written in matrix format, i.e., $y = Ax$ for some $m \times n$ matrix




- Example: $Ax = \begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \\ a_{31} & a_{32} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} a_{11}x_1 + a_{12}x_2 \\ a_{21}x_1 + a_{22}x_2 \\ a_{31}x_1 + a_{32}x_2 \end{bmatrix}$

- Example: $Ax = [\vec{a}_1 \ \vec{a}_2] \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = x_1\vec{a}_1 + x_2\vec{a}_2$, where \vec{a}_i is i -th column of A

– Output is linear combination of columns of A

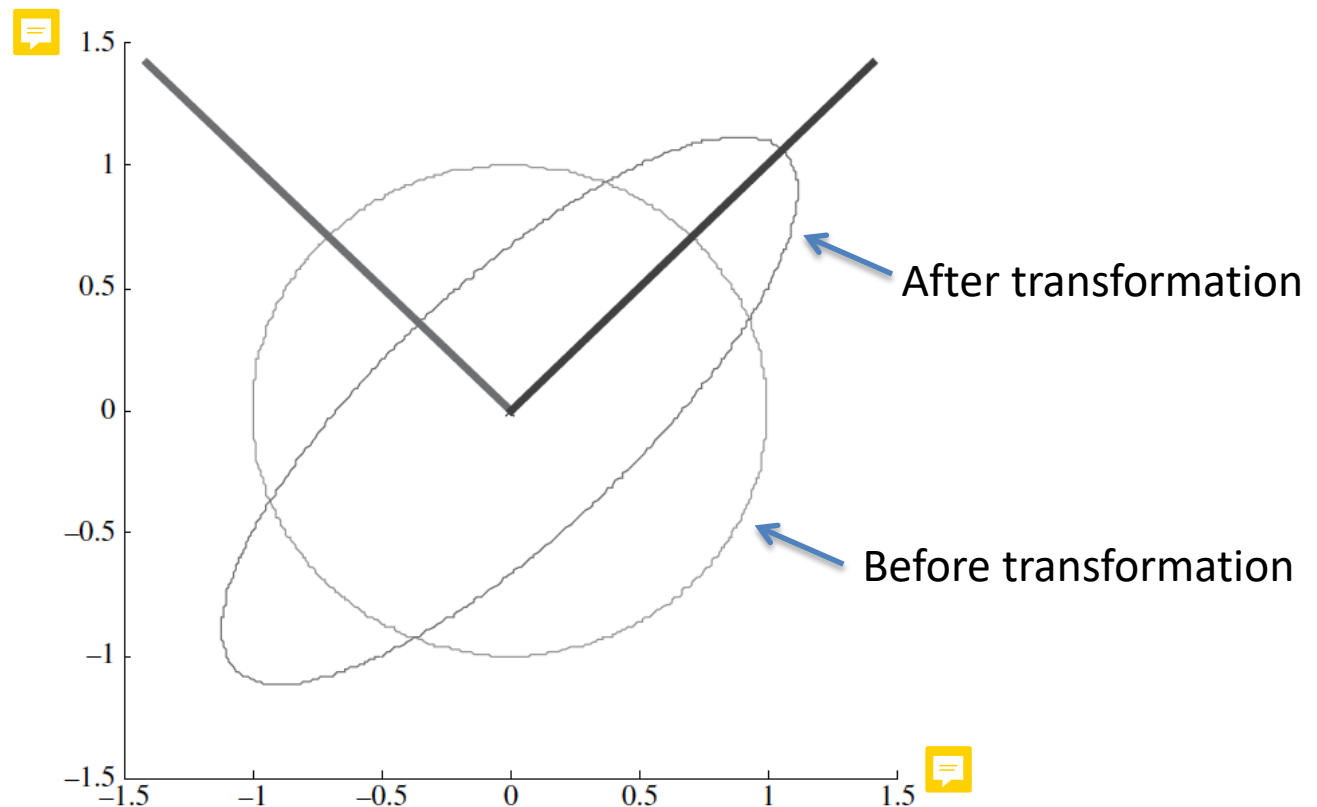
Linear Transformation: Example 1

- $A = \begin{bmatrix} -1 & 0 \\ 0 & 1 \end{bmatrix}$ flips vector about vertical axis




Linear Transformation: Example 2

- $A = \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix} \Rightarrow \begin{bmatrix} 1 & 0.5 \\ 0.5 & 1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} x_1 + 0.5x_2 \\ 0.5x_1 + x_2 \end{bmatrix}$
- **Circle** of points (below) mapped by A onto an **ellipse** (below)




Eigenvectors & Eigenvalues

- Consider linear system $y = Ax$, where $A = M \times M$ matrix
- Suppose x is eigenvector of A
 - By definition of eigenvectors, $Ax = \lambda x$ (for some number λ)
 - Therefore **direction of x preserved** under transformation A




$x = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$ is eigenvector with $\lambda = 1$

$\begin{bmatrix} \frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}$ & $\begin{bmatrix} -\frac{1}{\sqrt{2}} \\ \frac{1}{\sqrt{2}} \end{bmatrix}$ eigenvectors with eigenvalues 1.5 and 0.5



Computing Eigenvectors/Eigenvalues

- $u \neq 0$ is an eigenvector of $M \times M$ matrix A if $Au = \lambda u$ for some λ 

$$Au = \lambda u$$

$$(A - \lambda I)u = 0$$


$$\det(A - \lambda I) = 0$$



u is non-zero $\Rightarrow A - \lambda I$ is not full rank

- For example, $A = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$, then

$$\det \begin{pmatrix} a - \lambda & b \\ c & d - \lambda \end{pmatrix} = 0$$

 $(a - \lambda)(d - \lambda) - bc = 0$

$$\lambda^2 - (a + d)\lambda + (ad - bc) = 0$$

- Quadratic equations \Rightarrow up to two λ s (in general up to M eigenvalues for $M \times M$ matrix A)
- Solve for λ s, and then eigenvectors
- In real applications, use software to solve (e.g., 'eig' function in matlab)

Matrix Eigen-Decomposition

- For $M \times M$ matrix A with M eigenvalues $\lambda_1, \dots, \lambda_M$ and corresponding eigenvectors u_1, \dots, u_M , where $u_m^T u_m = 1$ (unit length)

$$\begin{aligned}
 A[u_1 \ u_2 \ \cdots \ u_M] &= [\lambda_1 u_1 \ \lambda_2 u_2 \ \cdots \ \lambda_M u_M] \\
 &= [u_1 \ u_2 \ \cdots \ u_M] \begin{bmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_M \end{bmatrix} \\
 AU &= U\Lambda,
 \end{aligned}$$

where Λ is a diagonal matrix which is all 0 except the diagonals which correspond to eigenvalues and U is a matrix whose columns are eigenvectors



- A symmetric $\implies M$ (not necessarily unique) eigenvalues, and M orthogonal eigenvectors, i.e., $u_i^T u_j = \delta(i - j) \implies U^{-1} = U^T$

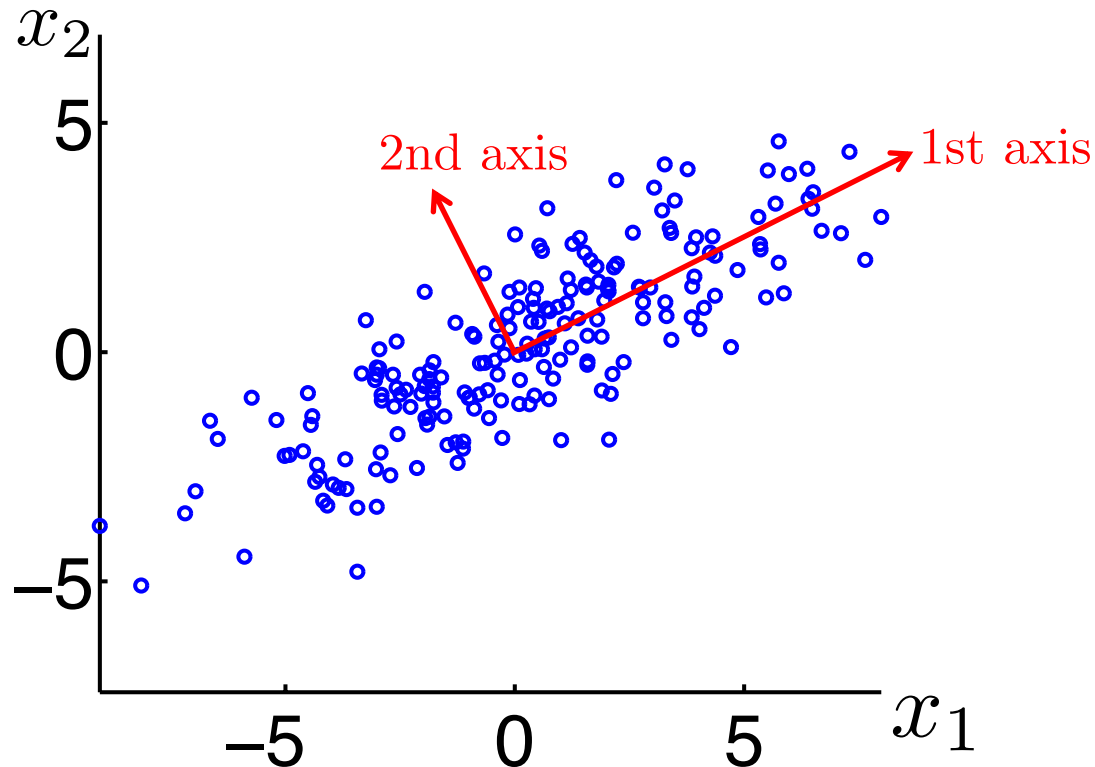
$$AU = U\Lambda \implies A = U\Lambda U^T = \lambda_1 u_1 u_1^T + \lambda_2 u_2 u_2^T \cdots + \lambda_M u_M u_M^T = \sum_m \lambda_m u_m u_m^T$$

- A positive semidefinite, i.e., $x^T A x \geq 0$ for all $x \implies \lambda_m \geq 0$ for all m .

Principal Component Analysis (PCA)


Principal Component Analysis (PCA)

- PCA = new coordinate system with orthogonal axes
 - Assume mean of input data = $\vec{0}$
 - 1st axis direction of largest variability in data 
 - 2nd axis direction of largest variability **orthogonal** to 1st axis 
 - 3rd axis direction of largest variability **orthogonal** to 1st and 2nd axes
 - ...



Projecting Onto Principal Axes

- Coordinates with respect to new axes given by projection onto axes
 - For D -dimensional point x , new coordinates = $(v_1^T x, v_2^T x, \dots, v_K^T x)$, where v_i is unit vector in same direction as i -th axis
 - x now K -dimensional vector: we have achieved dimensionality reduction for $K < D$
 - $v_i^T x$ called i -th principal component (PC)



1st Principal Axis is Direction of Largest Variability

- Consider $x^{(1)}, \dots, x^{(N)}$ (N data points, each dimension D)
- Assume data centered at origin, i.e., $\frac{1}{N} \sum_n x^{(n)} = \vec{0}$
- Let v_1 = unit vector in same direction as 1st principal axis, so resulting projected data are $[v_1^T x^{(1)}, \dots, v_1^T x^{(N)}]$
- Note that projected data has $\frac{1}{N} \sum_n v_1^T x^{(n)} = 0$ mean, so (empirical) variance of projected data are

$$\frac{1}{N} \sum_{n=1}^N \left(v_1^T x^{(n)} \right)^2 = \frac{1}{N} \sum_{n=1}^N v_1^T x^{(n)} x^{(n)T} v_1 = v_1^T \left(\frac{1}{N} \sum_{n=1}^N x^{(n)} x^{(n)T} \right) v_1 = v_1^T \Sigma v_1$$

where Σ is covariance of data samples

$$- \Sigma = \frac{1}{N} X X^T, \text{ where } X \text{ is the } D \times N \text{ matrix } [x^{(1)}, \dots, x^{(N)}]$$

- Therefore 1st principal axis (direction of data with largest variance):

$$v_1 = \underset{v}{\operatorname{argmax}} v^T \Sigma v \quad \text{where} \quad v^T v = 1$$

Finding Principal Axes

- 1st principal axis (direction of data with largest variance):

$$v_1 = \operatorname{argmax}_v v^T \Sigma v \quad \text{where} \quad v^T v = 1$$

- Σ is symmetric & positive semidefinite, so has D eigenvalues $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_D \geq 0$ with eigenvectors u_1, \dots, u_D , where $u_i^T u_j = \delta(i - j)$:

$$\Sigma = \lambda_1 u_1 u_1^T + \lambda_2 u_2 u_2^T + \dots + \lambda_D u_D u_D^T$$

- Therefore

$$\begin{aligned} v^T \Sigma v &= v^T (\lambda_1 u_1 u_1^T + \lambda_2 u_2 u_2^T + \dots + \lambda_D u_D u_D^T) v \\ &= \lambda_1 (v^T u_1)^2 + \lambda_2 (v^T u_2)^2 + \dots + \lambda_D (v^T u_D)^2 \end{aligned}$$

- Since λ_1 largest and $v_1^T v_1 = 1$, therefore $v_1 = u_1$

- $v_2 = \operatorname{argmax}_v v^T \Sigma v$ where $v^T v = 1$ and $v_1^T v_2 = 0 \implies v_2 = u_2$

- In general, $v_n = u_n$

Principal Component Analysis (PCA)


- Consider $x^{(1)}, \dots, x^{(N)}$ (N data points, each dimension D)
- Compute mean $\mu = \frac{1}{N} \sum_n x^{(n)}$
- Subtract mean from each sample point $\bar{x}^{(n)} = x^{(n)} - \mu$ (this ensures resulting data are centered at the origin) – procedure called “demean”
- Form $D \times N$ matrix $X = [\bar{x}^{(1)}, \dots, \bar{x}^{(N)}]$
- Compute sample covariance matrix $\Sigma = \frac{1}{N} X X^T$
- Find top K eigenvalues and corresponding eigenvectors v_1, \dots, v_K
- Given new (or old) datapoint x , we can represent x by new coordinates $[v_1^T(x - \mu) \ v_2^T(x - \mu) \ \dots \ v_K^T(x - \mu)]$
- Why top K eigenvalues and not bottom K ?

Principal Component Analysis (PCA) Representation

What does PCA representation buy us?

- Consider $x^{(1)}, \dots, x^{(N)}$ (N data points, each dimension D) with 0 mean
- Perform PCA and represent **old** datapoint $x^{(n)}$ by $[v_1^T x^{(n)} \quad v_2^T x^{(n)} \quad \dots \quad v_K^T x^{(n)}]$
- Can “reconstruct” $x^{(n)}$ by $\hat{x}^{(n)} = \sum_{k=1}^K (v_k^T x^{(n)}) v_k$
- Consider reconstruction error

$$\|x^{(n)} - \hat{x}^{(n)}\|^2 \triangleq \sum_{i=1}^D (x_i^{(n)} - \hat{x}_i^{(n)})^2$$



- Total reconstruction error (proof at end of notes):

$$\sum_{n=1}^N \|x^{(n)} - \hat{x}^{(n)}\|^2 = \sum_{n=1}^N x^{(n)T} x^{(n)} - N \sum_{i=1}^K v_i^T \Sigma v_i$$

- First term in reconstruction error always positive \implies to minimize reconstruction error, want second term $\sum_{i=1}^K v_i^T \Sigma v_i$ as big as possible, which is how PCA was defined \implies PCA lets us represent our original data with the least amount of error with K coefficients

Remarks

- Previous argument still works when $x^{(1)}, \dots, x^{(N)}$ has non-zero mean μ
- Just need to subtract μ from data before PCA
- When projecting datapoint x onto principal axes, new coordinates are $[v_1^T(x - \mu), v_2^T(x - \mu), \dots, v_K^T(x - \mu)]$
- Reconstructed $\hat{x} = \mu + \sum_{k=1}^K (v_k^T(x - \mu)) v_k$
- No magic way to select K . One way is the following criterion

$$\frac{\sum_{i=1}^K \lambda_i}{\sum_{i=1}^D \lambda_i} \geq \text{threshold (e.g., 0.95)}$$

- For $M \times M$ images, can convert into a vector (i.e., $D = M^2$) by stacking columns of pixel values into one single column


Eigenfaces Example

Eigenfaces

- PCA on 200 face images



16 of 200 faces




Mean of 200 faces

Eigenfaces


- PCA on 200 face images



16 of 200 faces
after de-meaning



Mean of 200 faces
(origin of coordinate
System)



Top 9 Eigenvectors
(Eigenfaces)

See `eigenfaces.m` on IVLE

Eigenfaces

- PCA on 200 face images



16 of 200 faces




Reconstructed Faces
From Top 50 eigenfaces

See `eigenfaces.m` on IVLE

Summary

- Curse of dimensionality: pattern recognition harder for high dimensions
 \implies useful to reduce dimensions
- PCA finds orthonormal basis (coordinate system) for data
- Principal axes corresponds to eigenvectors of data covariance matrix
- Principal axes sorted in order of importance (based on eigenvalues)
- Can discard axes with lower eigenvalues (hoping they are noise or unimportant for application)
- Reduce data dimensions by projecting onto remaining axes
- PCA optimal in the sense that it minimizes reconstruction error as measured by Euclidean distance
- PCA does not work for nonlinear space



Optional Reading


- Duda and Hart: Pattern Recognition, Chapter 3.8.1

Additional Material: PCA minimizes
reconstruction error proof

What does PCA representation buy us?

- Consider $x^{(1)}, \dots, x^{(N)}$ (N data points, each dimension D) with 0 mean
- Perform PCA and represent **old** datapoint $x^{(n)}$ by $[v_1^T x^{(n)} \ v_2^T x^{(n)} \ \dots \ v_K^T x^{(n)}]$
- Can “reconstruct” $x^{(n)}$ by $\hat{x}^{(n)} = \sum_{k=1}^K (v_k^T x^{(n)}) v_k$
- Consider reconstruction error

$$\begin{aligned}
 \|x^{(n)} - \hat{x}^{(n)}\|^2 &\triangleq \sum_{i=1}^D (x_i^{(n)} - \hat{x}_i^{(n)})^2 \\
 &= (x^{(n)} - \hat{x}^{(n)})^T (x^{(n)} - \hat{x}^{(n)}) \\
 &= \left(x^{(n)} - \sum_{k=1}^K (v_k^T x^{(n)}) v_k \right)^T \left(x^{(n)} - \sum_{k=1}^K (v_k^T x^{(n)}) v_k \right) \\
 &= x^{(n)T} x^{(n)} - 2 \sum_{k=1}^K (v_k^T x^{(n)})^2 + \sum_{k=1}^K (v_k^T x^{(n)})^2 \\
 &= x^{(n)T} x^{(n)} - \sum_{k=1}^K (v_k^T x^{(n)})^2
 \end{aligned}$$



What does PCA representation buy us?

- From previous slide: $\|x^{(n)} - \hat{x}^{(n)}\|^2 = x^{(n)T}x^{(n)} - \sum_{i=1}^K (v_i^T x^{(n)})^2$
- Total reconstruction error:

$$\begin{aligned}\sum_{n=1}^N \|x^{(n)} - \hat{x}^{(n)}\|^2 &= \sum_{n=1}^N x^{(n)T}x^{(n)} - \sum_{n=1}^N \sum_{i=1}^K (v_i^T x^{(n)})^2 \\&= \sum_{n=1}^N x^{(n)T}x^{(n)} - \sum_{n=1}^N \sum_{i=1}^K v_i^T x^{(n)} x^{(n)T} v_i \\&= \sum_{n=1}^N x^{(n)T}x^{(n)} - \sum_{i=1}^K v_i^T \left(\sum_{n=1}^N x^{(n)} x^{(n)T} \right) v_i \\&= \sum_{n=1}^N x^{(n)T}x^{(n)} - N \sum_{i=1}^K v_i^T \Sigma v_i\end{aligned}$$

- First term in reconstruction error always positive \implies to minimize reconstruction error, want second term $\sum_{i=1}^K v_i^T \Sigma v_i$ as big as possible, which is how PCA was defined \implies PCA lets us represent our original data with the least amount of error with K coefficients